



Replica scaling

Qwen3-8B BF16 | RTX 5090 | vLLM | 2,048 input / 256 output | TP=1 | 3 repeats per configuration


Output throughput



Mean run P95 TTFT



Mean run P95 TPOT



Peak scaling efficiency

