



Prefix caching

Qwen3-8B BF16 | RTX 5090 | vLLM | 4 GPUs | TP=1 | concurrency 32 | prefix cache enabled


Mean run P95 TTFT



Mean run P95 prefill latency



Mean run P95 TPOT



Resident prefix cache

