



Tensor parallelism

Qwen3-8B BF16 | RTX 5090 | vLLM | balanced workload | fixed 8-GPU budget | replicas = 8 / TP


Output throughput at concurrency 32



Mean run P95 TTFT at concurrency 1



GPU memory at concurrency 32



Collective time at concurrency 32

