



Workload profiles

Qwen3-8B BF16 | RTX 5090 | vLLM | 4 GPUs | TP=1 | five request-length profiles


Output throughput



Joint SLO goodput



Mean run P95 TTFT



Mean run P95 end-to-end latency

