



# Request routing

Qwen3-8B BF16 | RTX 5090 | vLLM | RAG | 8 GPUs | concurrency 128 | prefix reuse 50%


Mean run P95 TTFT



Mean run P95 queue latency



Worker load imbalance index



Joint SLO goodput

